Natural Language Inference

By Jack Pay & Dennis Kiselev

Introduction

- **Two models** were developed to solve Natural Language Inference (**NLI**), the task of assessing whether a premise semantically supports a hypothesis
- Method B focused on using RNNs whilst C finetuned pretrained Transformer architectures

Methods

- **Method B**:
  - Trained a **Bi-LSTM** with frozen **XLNET** embeddings
  - Utilised **learning rate scheduling** to approach a global optimum
    - LR was reduced on metric plateau
  - Employed **subtractive & multiplicative sentence fusion**, & **attention** to enhance sentence representations
- **Method C**:
  - Finetuned the base **ROBERTA** Transformer model
  - Employed **data augmentation** to bolster the training data:
    - Synonym replacements & insertions
    - Word deletions & swaps
  - Used **early stopping** and LR scheduling to reduce overfitting
    - LR warmup and then decreasing was used

Conclusions

- **Pretrained word embeddings** are key to higher performance
- **Transformers** achieve a **higher performance** compared with RNN
- Basic textual augmentations do not significantly improve performance
- Models struggle when premise & hypotheses are **ambiguous**, contain **disfluencies** or **OOV** words

References

- S. Zhang, S. Liu and M. Liu, "Natural language inference using LSTM model with sentence fusion," 2017 36th Chinese Control Conference (CCC), Dalian, China, 2017, pp. 11081-11085, doi: 10.23919/ChiCC.2017.8029126
- Tarunesh, Ishan, Somak Aditya, and Monojit Choudhury. "Trusting roberta over bert: Insights from checklisting the natural language inference task." arXiv preprint arXiv:2107.07229 (2021).
- Sadat, Mobashir, and Cornelia Caragea. "Scinli: A corpus for natural language inference on scientific text." arXiv preprint arXiv:2203.06728 (2022).

Task B: Figures

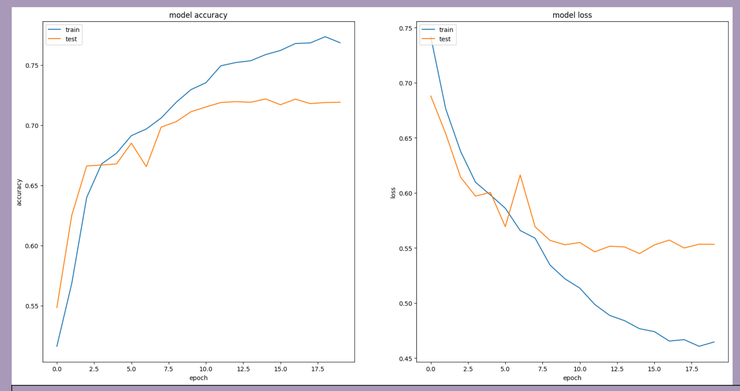

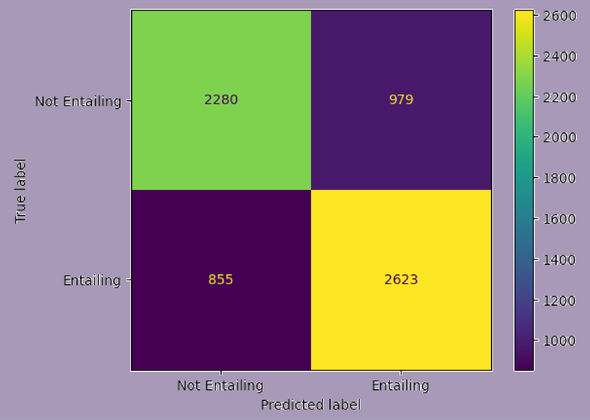

Predictions with the Highest Loss

Loss	Premise	Hypothesis
11.88	He was the first scholar to describe in detail.	He was the first scholar to describe in detail and document the long-term cyclical relationship between global population cycles and cycles of political rebellion and revolution.
6.59	La Barredora ("The Sweeper Truck") is a criminal gang based in the Mexican resort city of Acapulco, Guerrero and its surrounding territories.	La Barredora ("The Sweeper Truck") is a criminal gang based in the Mexican resort city of Acapulco, City and County of Denver and its surrounding territories.
5.58	Traditional examples of background music include music played at various social gatherings and music played in certain retail venues.	Traditional examples of background music include music played at same social gatherings and music played in certain retail venues.

Task C: Figures

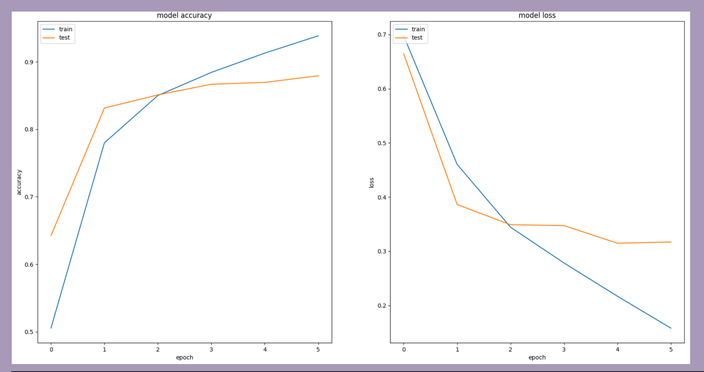

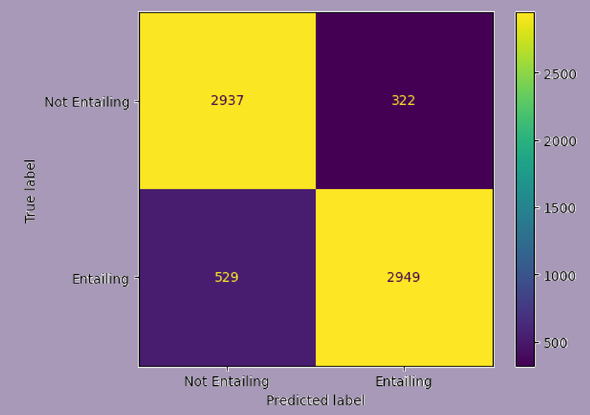

Predictions with the Highest Loss

Loss	Premise	Hypothesis
6.04	okay well what are we doing about recycling in our community	Our community does not recycle.
5.90	The port is virtually round the corner from the passeig.	It takes at least two days to go from the port to the passeig.
5.87	and i don't know about your part of the country but uh down here in the last year oh year plus i it was last beginning with last year's Earth Day	I've never been to your part of the country.

Optimal Results

Method	Accuracy	Precision	Macro Precision	Weighted Macro Precision	Recall	Macro Recall	Weighted Macro Recall	F1-Score	Macro F1-Score	Weighted Macro F1-Score	MCC	Loss
B	0.728	0.728	0.728	0.728	0.754	0.727	0.728	0.741	0.727	0.728	0.455	0.535
C	0.874	0.902	0.874	0.875	0.848	0.875	0.874	0.874	0.874	0.874	0.749	0.333